

Improved Classical and Quantum Algorithms for the Shortest Vector Problem

Divesh Aggarwal



Yanlin Chen



Rajendra Kumar




Yixin Shen




What is a (Euclidean) lattice?

Definition

$\mathcal{L}(\mathbf{b}_1, \dots, \mathbf{b}_n) = \left\{ \sum_{i=1}^n x_i \mathbf{b}_i : x_i \in \mathbb{Z} \right\}$ where $\mathbf{b}_1, \dots, \mathbf{b}_n$ is a basis of \mathbb{R}^n .




Lattice-based cryptography: fundamental idea



- ▶ **good basis**: private information, makes problem easy
- ▶ **bad basis**: public information, makes problem hard

Lattice-based cryptography: fundamental idea




- ▶ **good basis**: private information, makes problem easy
- ▶ **bad basis**: public information, makes problem hard

Basis reduction: transform a bad basis into a good one

Main tool: BKZ algorithm and its variants

Requires to solve the **(approx-)SVP problem** in smaller dimensions.

The Shortest Vector Problem




Shortest Vector Problem (SVP):

Given a basis for the lattice \mathcal{L} , find a shortest nonzero lattice vector.

$\lambda_1(\mathcal{L}) = \text{length of such a vector.}$

The Shortest Vector Problem



Shortest Vector Problem (SVP):

Given a basis for the lattice \mathcal{L} , find a shortest nonzero lattice vector.

$\lambda_1(\mathcal{L}) = \text{length of such a vector.}$

Main approaches for SVP:

- ▶ Enumeration: $2^{O(n \log(n))}$ time and $\text{poly}(n)$ space
- ▶ Sieving: $2^{O(n)}$ time and $2^{O(n)}$ space

Sieving

- ▶ Heuristic algorithms: fastest in practice
- ▶ Provable algorithms: important for theory → our work

Results in the Classical Setting

Provable algorithms for SVP:

Time Complexity	Space Complexity	Reference
$n^{\frac{n}{2e}+o(n)}$	$\text{poly}(n)$	[Kan87,HS07]
$2^{n+o(n)}$	$2^{n+o(n)}$	[ADRS15]
$2^{2.05n+o(n)}$	$2^{0.5n+o(n)}$	[CCL18]
$2^{1.669n+o(n)}$	$2^{0.5n+o(n)}$	Our work

Results in the Classical Setting

Provable algorithms for SVP:

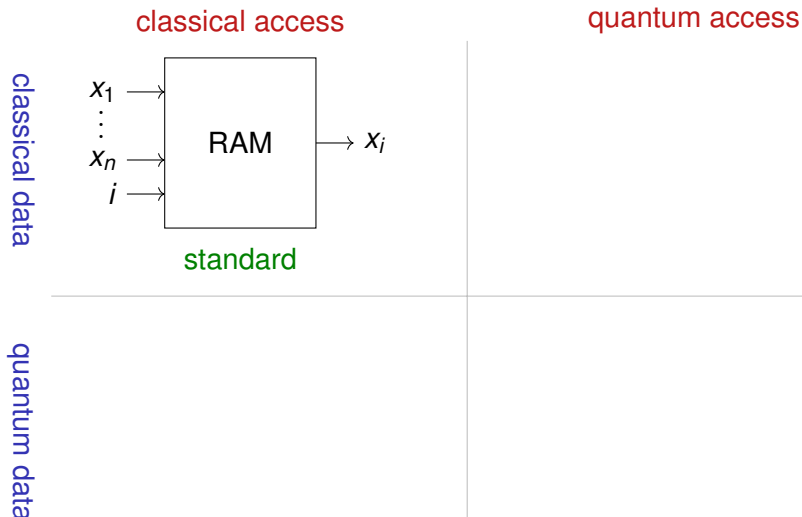
Time Complexity	Space Complexity	Reference
$n^{\frac{n}{2e}+o(n)}$	$\text{poly}(n)$	[Kan87,HS07]
$2^{n+o(n)}$	$2^{n+o(n)}$	[ADRS15]
$2^{2.05n+o(n)}$	$2^{0.5n+o(n)}$	[CCL18]
$2^{1.669n+o(n)}$	$2^{0.5n+o(n)}$	Our work

Our work: first provable smooth time/space trade-off for SVP

$$\text{time } q^{13n+o(n)} \quad \text{space } \text{poly}(n) \cdot q^{\frac{16n}{q^2}} \quad q \in [4, \sqrt{n}]$$

- ▶ $q = \sqrt{n}$: time $n^{O(n)}$ and space $\text{poly}(n)$, not as good as [Kan87].
- ▶ $q = 4$: time $2^{O(n)}$ and space $2^{O(n)}$, not as good as [ADRS15].

Interlude: quantum memory models




Assumption: $O(1)$ time cost

Interlude: quantum memory models

classical access


quantum access

classical data



standard

quantum data




standard

Assumption: $O(1)$ time cost


Interlude: quantum memory models

classical access




standard

quantum access



potentially strong assumption

quantum data




standard

Assumption: $O(1)$ time cost


Interlude: quantum memory models

classical access



standard


quantum access




potentially strong assumption

classical data

quantum data



standard



strong assumption

Assumption: $O(1)$ time cost

Results in the Quantum Setting

Provable quantum algorithms for SVP:

Time Complexity	Space Complexity			Reference
	Classical	Quantum	Model	
$2^{1.799n+o(n)}$	$2^{1.286n+o(n)}$	$2^{1.286n+o(n)}$	QRACM	[LMP15]
$2^{1.2553n+o(n)}$	$2^{0.5n+o(n)}$	$\text{poly}(n)$	plain	[CCL18]
$2^{n+o(n)}$	$2^{n+o(n)}$	classical algorithm!		[ADRS15]
$2^{0.950n+o(n)}$	$2^{0.5n+o(n)}$	$\text{poly}(n)$	plain	Our work
$2^{0.835n+o(n)}$	$2^{0.5n+o(n)}$	$2^{0.293n+o(n)}$	QRACM	Our work

Remark on quantum heuristic algorithms:

- ▶ better complexity: $2^{0.265n+o(n)}$ [Laarhoven15], requires QRACM
- ▶ even better complexity: $2^{0.257n+o(n)}$ [CL21], requires QRAQM

Sieving Algorithms

Original idea [AKS01]:

- ▶ Reduce basis
- ▶ Generate random vectors
- ▶ Repeat many times:
 - ▶ Sieve vectors

Sieving Algorithms

Original idea [AKS01]:

- ▶ Reduce basis
- ▶ Generate random vectors
- ▶ Repeat many times:
 - ▶ Sieve vectors

Sieve:

Input: many vectors of length $\leq \ell$

Output: many vectors of length $\leq \frac{\ell}{2}$

Combine pairs of vectors to produce shorter vectors

Sieving Algorithms

Original idea [AKS01]:

- ▶ Reduce basis
- ▶ Generate random vectors
- ▶ Repeat many times:
 - ▶ Sieve vectors

Sieve:

Input: many vectors of length $\leq \ell$

Output: many vectors of length $\leq \frac{\ell}{2}$

Combine pairs of vectors to produce shorter vectors

Idea: LLL reduced \leadsto length $\ell \leq 2^{O(n)} \lambda_1$, sieve $O(n)$, solve SVP

Sieving Algorithms

Original idea [AKS01]:

- ▶ Reduce basis
- ▶ Generate random vectors
- ▶ Repeat many times:
 - ▶ Sieve vectors

Sieve:

Input: many vectors of length $\leq \ell$

Output: many vectors of length $\leq \frac{\ell}{2}$

Combine pairs of vectors to produce shorter vectors

Idea: LLL reduced \leadsto length $\ell \leq 2^{O(n)} \lambda_1$, sieve $O(n)$, solve SVP

Many heuristic variants: local sensitive hash, tuple sieve, ...

All control the **length** of the vectors.

Sieving Algorithms

Original idea [AKS01]:

- ▶ Reduce basis
- ▶ Generate random vectors
- ▶ Repeat many times:
 - ▶ Sieve vectors

Sieve:

Input: many vectors of length $\leq \ell$

Output: many vectors of length $\leq \frac{\ell}{2}$

Combine pairs of vectors to produce shorter vectors

Idea: LLL reduced \leadsto length $\ell \leq 2^{O(n)} \lambda_1$, sieve $O(n)$, solve SVP

Many heuristic variants: local sensitive hash, tuple sieve, ...

All control the **length** of the vectors.


[ADRS15]'s new idea: control **distribution** instead of length of vectors

Discrete Gaussian Sampling


$$\rho_s(\mathbf{x}) = \exp\left(-\pi \frac{\|\mathbf{x}\|^2}{s^2}\right), \quad D_{L,s}(\mathbf{x}) = \frac{\rho_s(\mathbf{x})}{\rho_s(L)}, \quad \mathbf{x} \in \mathbb{R}^n, s > 0.$$

Definition (Discrete Gaussian Distribution)


On lattice L with **parameter** s : probability of $\mathbf{x} \in L$ is $D_{L,s}(\mathbf{x})$.



$$L = \mathbb{Z}, s = 7$$



$$L = \mathbb{Z}^2, s = 7$$



$$L = \mathbb{Z} \times 4\mathbb{Z}, s = 7$$



Discrete Gaussian Sampling (DGS)

- ▶ **input:** L and s
- ▶ **output:** random $\mathbf{x} \in L$ according to $D_{L,s}$.

DGS, BDD and SVP





DGS, BDD and SVP



Bounded Distance Decoding (α -BDD):

Given a lattice \mathcal{L} and a target vector $t \in \mathbb{R}^n$



DGS, BDD and SVP



Bounded Distance Decoding (α -BDD):

Given a lattice \mathcal{L} and a target vector $t \in \mathbb{R}^n$ with distance to lattice $\leq \alpha \cdot \lambda_1(\mathcal{L})$

DGS, BDD and SVP





Bounded Distance Decoding (α -BDD):

Given a lattice \mathcal{L} and a target vector $t \in \mathbb{R}^n$ with distance to lattice $\leq \alpha \cdot \lambda_1(\mathcal{L})$, find the closest vector $y \in \mathcal{L}$.

- ▶ α is the decoding radius
- ▶ $\alpha < \frac{1}{2}$ for unique solution

DGS, BDD and SVP



Bounded Distance Decoding (α -BDD):


Given a lattice \mathcal{L} and a target vector $t \in \mathbb{R}^n$ with distance to lattice $\leq \alpha \cdot \lambda_1(\mathcal{L})$, find the closest vector $y \in \mathcal{L}$.

- ▶ α is the decoding radius
- ▶ $\alpha < \frac{1}{2}$ for unique solution

The two reductions use completely different DGS parameter regimes!


Hardness of Discrete Gaussian Sampling

Parameter s (width/standard deviation) of $D_{\mathcal{L},s}$:




Hardness of Discrete Gaussian Sampling

Parameter s (width/standard deviation) of $D_{\mathcal{L},s}$:




Hardness of Discrete Gaussian Sampling

Parameter s (width/standard deviation) of $D_{\mathcal{L},s}$:



Hardness of Discrete Gaussian Sampling

Parameter s (width/standard deviation) of $D_{\mathcal{L},s}$:




Theorem (ADRS15, best known result)


There is an algorithm that solves DGS for $s = \sqrt{2}\eta_{1/2}(\mathcal{L})$ in time and space $2^{n/2+o(n)}$.

Hardness of Discrete Gaussian Sampling

Parameter s (width/standard deviation) of $D_{\mathcal{L},s}$:



- ▶ hard to sample
- ▶ SVP
- smoothing parameter
- ▶ easy to sample
- ▶ Open problem: $2^{O(n)}$ time, $2^{o(n)}$ space algorithm for $s = \eta_\epsilon(\mathcal{L})$
- ▶ No known time/space trade-off for $s \ll \eta_\epsilon(\mathcal{L})$



\leadsto first provable time/space trade-off for SVP

DGS time/space trade-off (simplified)

Idea: if $X_1, \dots, X_k \sim D_{\mathcal{L}, s}$ and $\sum_i X_i \in q\mathcal{L}$ **then** $(\sum_i X_i)/q \approx D_{\mathcal{L}, s\sqrt{k}/q}$
 \leadsto progress when $k < q^2$, repeat many times to reach $\eta_\epsilon(\mathcal{L})$

DGS time/space trade-off (simplified)

Idea: if $X_1, \dots, X_k \sim D_{\mathcal{L}, \mathbf{s}}$ and $\sum_i X_i \in q\mathcal{L}$ **then** $(\sum_i X_i)/q \approx D_{\mathcal{L}, \mathbf{s}\sqrt{k}/q}$
 \leadsto progress when $k < q^2$, repeat many times to reach $\eta_\epsilon(\mathcal{L})$

Algorithm: given a list of N vectors in \mathcal{L} , find $k = q^2 - 1$ of them such that their sum $\in q\mathcal{L}$, then repeat (q is a parameter)

DGS time/space trade-off (simplified)

Idea: if $X_1, \dots, X_k \sim D_{\mathcal{L}, s}$ and $\sum_i X_i \in q\mathcal{L}$ **then** $(\sum_i X_i)/q \approx D_{\mathcal{L}, s\sqrt{k}/q}$
 \leadsto progress when $k < q^2$, repeat many times to reach $\eta_\epsilon(\mathcal{L})$

Algorithm: given a list of N vectors in \mathcal{L} , find $k = q^2 - 1$ of them such that their sum $\in q\mathcal{L}$, then repeat (q is a parameter)

- **Space:** need $N \gtrsim q^{n/q^2}$ to be successful
- **Time:** q^n to produce one vector

decrease with q
increase with q

DGS time/space trade-off (simplified)

Idea: if $X_1, \dots, X_k \sim D_{\mathcal{L}, s}$ and $\sum_i X_i \in q\mathcal{L}$ then $(\sum_i X_i)/q \approx D_{\mathcal{L}, s\sqrt{k}/q}$
 \leadsto progress when $k < q^2$, repeat many times to reach $\eta_\epsilon(\mathcal{L})$

Algorithm: given a list of N vectors in \mathcal{L} , find $k = q^2 - 1$ of them such that their sum $\in q\mathcal{L}$, then repeat (q is a parameter)

- ▶ **Space:** need $N \gtrsim q^{n/q^2}$ to be successful
 - ▶ **Time:** q^n to produce one vector
- decrease with q
increase with q


Difficulties:

- ▶ independence of samples
- ▶ errors in distributions

Theorem (Simplified)

For $q \in [4, \sqrt{n}]$, there is an algorithm that produces q^{16n/q^2} vectors from $D_{\mathcal{L}, s}$ with $s \geq \eta_\epsilon(\mathcal{L})$ in time q^{13n} and space q^{16n/q^2} .

Time-Space Tradeoff for SVP



Smooth time-space tradeoff for BDD: create a $\frac{0.1}{q}$ -BDD oracle in time q^{13n} , space q^{16n/q^2} , each call takes time q^{16n/q^2} .

Gives a smooth time-space tradeoff for SVP:

Theorem

Let $n \in \mathbb{N}$, $q \in [4, \sqrt{n}]$ be a positive integer. Let \mathcal{L} be a lattice of rank n . There is a randomized algorithm that solves SVP in time $q^{13n+o(n)}$ and in space $\text{poly}(n) \cdot q^{\frac{16n}{q^2}}$.

SVP to BDD reduction [CCL18]

Lemma (CCL18, simplified)

Given a α -BDD oracle and p an integer, one can enumerate all lattice points in a ball of radius $p\alpha\lambda_1$ using p^n queries to the oracle.

SVP to BDD reduction [CCL18]

Lemma (CCL18, simplified)

Given a α -BDD oracle and p an integer, one can enumerate all lattice points in a ball of radius $p\alpha\lambda_1$ using p^n queries to the oracle.

Solve SVP by using a α -BDD oracle:

- ▶ Set $p = \lceil \frac{1}{\alpha} \rceil$.
- ▶ Enumerate all points in a ball of radius $> \lambda_1$.

SVP to BDD reduction [CCL18]

Lemma (CCL18, simplified)

Given a α -BDD oracle and p an integer, one can enumerate all lattice points in a ball of radius $p\alpha\lambda_1$ using p^n queries to the oracle.

Solve SVP by using a α -BDD oracle:


- ▶ Set $p = \lceil \frac{1}{\alpha} \rceil$.
- ▶ Enumerate all points in a ball of radius $> \lambda_1$.

The reduction is space efficient

But $\alpha < \frac{1}{2} \implies p \geq 3 \implies$ at least 3^n queries

Faster SVP to BDD reduction


Cover the sphere of radius $\lambda_1(\mathcal{L})$ by balls of radius $2\alpha\lambda_1(\mathcal{L})$:



Use $2^n \alpha$ -BDD queries to enumerate points in balls of radius $2\alpha\lambda_1$

Faster SVP to BDD reduction

Cover the sphere of radius $\lambda_1(\mathcal{L})$ by balls of radius $2\alpha\lambda_1(\mathcal{L})$:




Use $2^n \alpha$ -BDD queries to enumerate points in balls of radius $2\alpha\lambda_1$

Each ball covers a spherical cap.

Faster SVP to BDD reduction

Cover the sphere of radius $\lambda_1(\mathcal{L})$ by balls of radius $2\alpha\lambda_1(\mathcal{L})$:



Use $2^n \alpha$ -BDD queries to enumerate points in balls of radius $2\alpha\lambda_1$

Each ball covers a spherical cap.


Smaller α :

- ▶ More balls
- ▶ Less expensive BDD

↪ Trade-off


SVP to DGS via BDD

Classical SVP to BDD: do 3^n queries to 1/3-BDD and keep minimum




SVP to DGS via BDD

Classical SVP to BDD: do 3^n queries to $1/3$ -BDD and keep minimum




“Improved” SVP to BDD: do $M(n, \alpha)$ queries to α -BDD

Details omitted in this presentation, M is a complicated function




SVP to DGS via BDD

Classical SVP to BDD: do 3^n queries to $1/3$ -BDD and keep minimum



“Improved” SVP to BDD: do $M(n, \alpha)$ queries to α -BDD

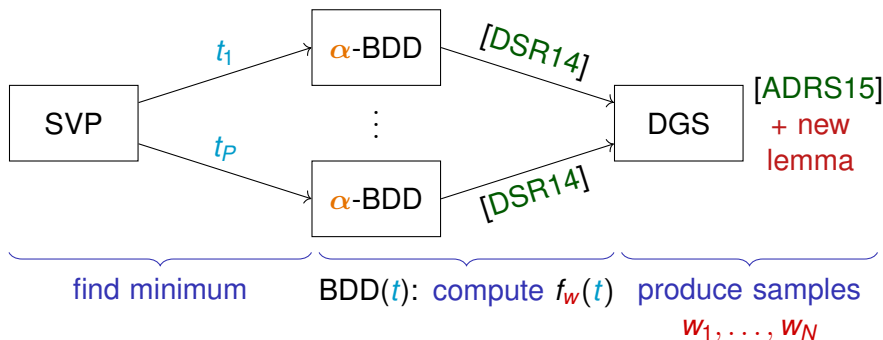
Details omitted in this presentation, M is a complicated function



- ▶ Not obvious which one is better: less queries to more expensive BDD oracle
- ▶ Same structure for both reductions, but different parameters:
will be useful for the quantum reduction!


Reduction

- ▶ SVP makes $P := M(n, \alpha)$ calls to α -BDD with argument t_1, \dots, t_P
- ▶ each BDD call requires N samples w_1, \dots, w_N from DGS
- ▶ w_1, \dots, w_N can be **shared** across all BDD calls: independent of t_i



Reduction

- ▶ SVP makes $P := M(n, \alpha)$ calls to α -BDD with argument t_1, \dots, t_P
- ▶ each BDD call requires N samples w_1, \dots, w_N from DGS
- ▶ w_1, \dots, w_N can be **shared** across all BDD calls: independent of t_i



Classical cost: DGS cost + $P \times$ BDD cost $\approx \text{poly}(n) \times (N + PN)$

Reduction from BDD to DGS

Periodic Gaussian function $f(\mathbf{t}) := \frac{\rho(\mathbf{t} + \mathcal{L})}{\rho(\mathcal{L})} = \mathbb{E}_{\mathbf{w} \sim \mathcal{D}_{\mathcal{L}^*}} [\cos(2\pi \langle \mathbf{w}, \mathbf{t} \rangle)]$

- ▶ f achieves maximum on lattice points
- ▶ a constant number of gradient ascent steps solves BDD

Reduction from BDD to DGS

Periodic Gaussian function $f(\mathbf{t}) := \frac{\rho(\mathbf{t} + \mathcal{L})}{\rho(\mathcal{L})} = \mathbb{E}_{\mathbf{w} \sim \mathcal{D}_{\mathcal{L}^*}} [\cos(2\pi \langle \mathbf{w}, \mathbf{t} \rangle)]$

- ▶ f achieves maximum on lattice points
- ▶ a constant number of gradient ascent steps solves BDD

Approximate f by

$$f_w(t) = \frac{1}{N} \sum_{i=1}^N \cos(2\pi \langle w_i, t \rangle)$$

where w_1, \dots, w_N are i.i.d. DGS samples: small error if N is very large.

Reduction from BDD to DGS

Periodic Gaussian function $f(\mathbf{t}) := \frac{\rho(\mathbf{t} + \mathcal{L})}{\rho(\mathcal{L})} = \mathbb{E}_{\mathbf{w} \sim \mathcal{D}_{\mathcal{L}^*}} [\cos(2\pi \langle \mathbf{w}, \mathbf{t} \rangle)]$

- ▶ f achieves maximum on lattice points
- ▶ a constant number of gradient ascent steps solves BDD

Approximate f by

$$f_w(t) = \frac{1}{N} \sum_{i=1}^N \cos(2\pi \langle w_i, t \rangle)$$

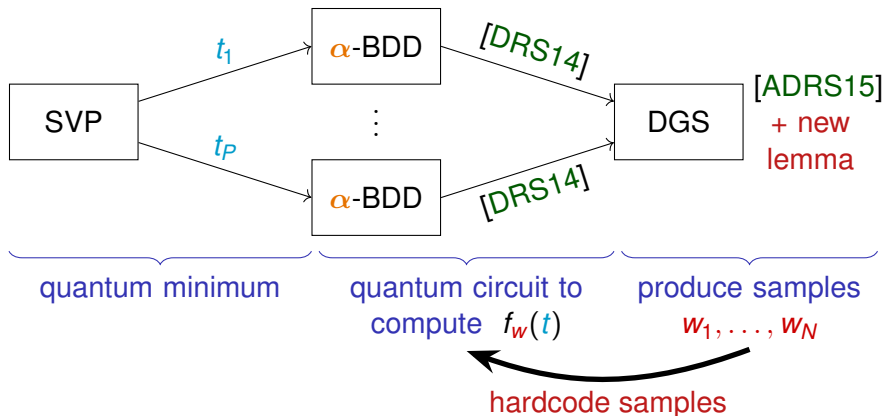
where w_1, \dots, w_N are i.i.d. DGS samples: small error if N is very large.

Theorem ([DRS14] (Informal))

There is an algorithm that solves α -BDD using N samples from $\mathcal{D}_{\mathcal{L}^, \eta_\varepsilon(\mathcal{L}^*)}$ in time $N \cdot \text{poly}(n)$, where $N = O\left(n^{\frac{\log(1/\varepsilon)}{\sqrt{\varepsilon}}}\right)$ and $\alpha = \alpha(\varepsilon)$.*


Quantum Reduction

- ▶ hardcode DGS samples into a quantum circuit to create a BDD oracle
- ▶ use this oracle in a quantum minimum finding algorithm



Quantum Reduction

- ▶ hardcode DGS samples into a quantum circuit to create a BDD oracle
- ▶ use this oracle in a quantum minimum finding algorithm



$$\text{Quantum cost: DGS cost} + \sqrt{P} \times \text{BDD cost} \approx \text{poly}(n) \times \left(N + \sqrt{PN} \right)$$

Reduction from BDD to DGS with QRACM

$$f_w(t) = \frac{1}{N} \sum_{i=1}^N \cos(2\pi \langle w_i, t \rangle)$$

where w_1, \dots, w_N are i.i.d. DGS samples: small error if N is very large.

Reduction from BDD to DGS with QRACM

$$f_w(t) = \frac{1}{N} \sum_{i=1}^N \cos(2\pi \langle w_i, t \rangle)$$

where w_1, \dots, w_N are i.i.d. DGS samples: small error if N is very large.

Our algorithm: approximate f_w quantumly in time $\sqrt{N} \cdot \text{poly}(n)$
Use amplitude estimation and show that the error stays small

Reduction from BDD to DGS with QRACM

$$f_w(t) = \frac{1}{N} \sum_{i=1}^N \cos(2\pi \langle w_i, t \rangle)$$

where w_1, \dots, w_N are i.i.d. DGS samples: small error if N is very large.

Our algorithm: approximate f_w quantumly in time $\sqrt{N} \cdot \text{poly}(n)$
Use amplitude estimation and show that the error stays small

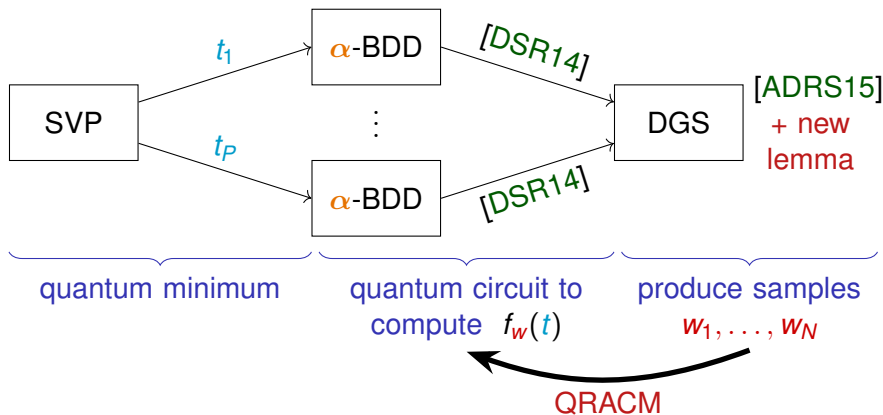
Theorem (Informal)

*There is an **quantum** algorithm that solves α -BDD using N samples from $D_{\mathcal{L}^*, \eta_\varepsilon(\mathcal{L}^*)}$ in time $\sqrt{N} \cdot \text{poly}(n)$, where $N = O\left(n^8 \frac{\log(1/\varepsilon)}{\sqrt{\varepsilon}}\right)$ and $\alpha = \alpha(\varepsilon)$. It requires a **QRACM of size N and $O(N)$ preprocessing time**.*

↪ gain when doing lots of BDD calls


Quantum Reduction with QRACM

- ▶ put DGS samples in a QRACM
- ▶ BDD oracle uses QRACM + amplitude estimation



Quantum Reduction with QRACM


- ▶ put DGS samples in a QRACM
- ▶ BDD oracle uses QRACM + amplitude estimation




Cost (QRACM): DGS cost + $\sqrt{P} \times \text{BDD cost} \approx \text{poly}(n) \times (N + \sqrt{PN})$

Quantum SVP

Classical SVP to BDD: do 3^n queries to 1/3-BDD and keep minimum



“Improved” SVP to BDD: do $M(n, \alpha)$ queries to α -BDD



- Not obvious which one is better: less queries to more expensive BDD oracle

SVP and Generalized Kissing Number

Number of lattice points in a ball of radius r is $\leq c^{n+o(n)} r^n$

$\beta(\mathcal{L})$ = smallest c that works for all r


- ▶ Upper bound: $\beta(\mathcal{L}) \leq 2^{0.401}$ [KL78]
- ▶ Conjectured to be $\beta(\mathcal{L}) \approx 1$ for most lattices

SVP and Generalized Kissing Number

Number of lattice points in a ball of radius r is $\leq c^{n+o(n)} r^n$

$\beta(\mathcal{L})$ = smallest c that works for all r

- ▶ Upper bound: $\beta(\mathcal{L}) \leq 2^{0.401}$ [KL78]
- ▶ Conjectured to be $\beta(\mathcal{L}) \approx 1$ for most lattices



Best known relations between α and ϵ depends on $\beta(\mathcal{L})$:


small $\beta(\mathcal{L}) \rightsquigarrow$ bigger α for fixed $\epsilon \rightsquigarrow$ less expensive BDD

SVP and Generalized Kissing Number

Number of lattice points in a ball of radius r is $\leq c^{n+o(n)} r^n$

$\beta(\mathcal{L})$ = smallest c that works for all r

- ▶ Upper bound: $\beta(\mathcal{L}) \leq 2^{0.401}$ [KL78]
- ▶ Conjectured to be $\beta(\mathcal{L}) \approx 1$ for most lattices




SVP and Generalized Kissing Number

Number of lattice points in a ball of radius r is $\leq c^{n+o(n)} r^n$

$\beta(\mathcal{L})$ = smallest c that works for all r

- ▶ Upper bound: $\beta(\mathcal{L}) \leq 2^{0.401}$ [KL78]
- ▶ Conjectured to be $\beta(\mathcal{L}) \approx 1$ for most lattices



Conclusions and Future work

Provable SVP:

- ▶ classical: time $2^{1.669n+o(n)}$, space $2^{0.5n+o(n)}$
- ▶ quantum: $2^{0.950n+o(n)}$, space $2^{0.5n+o(n)}$ and $\text{poly}(n)$ qubits
- ▶ quantum: $2^{0.835n+o(n)}$, classical space $2^{0.5n+o(n)}$ and QRACM $2^{0.293n+o(n)}$
- ▶ first time/space tradeoff: time q^{13n} , space q^{16n/q^2} for $q \in [4, \sqrt{n}]$
- ▶ studied dependency on $\beta(\mathcal{L})$, generalized kissing number

Conclusions and Future work

Provable SVP:

- ▶ classical: time $2^{1.669n+o(n)}$, space $2^{0.5n+o(n)}$
- ▶ quantum: $2^{0.950n+o(n)}$, space $2^{0.5n+o(n)}$ and $\text{poly}(n)$ qubits
- ▶ quantum: $2^{0.835n+o(n)}$, classical space $2^{0.5n+o(n)}$ and QRACM $2^{0.293n+o(n)}$
- ▶ first time/space tradeoff: time q^{13n} , space q^{16n/q^2} for $q \in [4, \sqrt{n}]$
- ▶ studied dependency on $\beta(\mathcal{L})$, generalized kissing number

Open problems:

- ▶ Show that random lattices satisfy $\beta(\mathcal{L}) \approx 1$?
- ▶ Fill the gap between provable and heuristic algorithms for sieving?
- ▶ Exploit the subexponential space regime in our trade-off for SVP?
- ▶ $2^{O(n)}$ time, $2^{o(n)}$ space algorithm for DGS at smoothing parameter?

Backup slides

DGS sampling: new lemma

- ▶ [ADRS15]: DGS of parameter $s \geq \sqrt{2}\eta_{1/2}(\mathcal{L})$ in time $2^{n/2}$
- ▶ BDD to DGS reduction requires $s = \eta_\varepsilon(\mathcal{L})$ for some $\varepsilon > 0$

Previous work [CCL18]: find ε such that $\eta_\varepsilon(\mathcal{L}) \geq \sqrt{2}\eta_{1/2}$

\leadsto very small ε , larger than necessary BDD radius, too expensive BDD

DGS sampling: new lemma

- ▶ [ADRS15]: DGS of parameter $s \geq \sqrt{2}\eta_{1/2}(\mathcal{L})$ in time $2^{n/2}$
- ▶ BDD to DGS reduction requires $s = \eta_\varepsilon(\mathcal{L})$ for some $\varepsilon > 0$

Previous work [CCL18]: find ε such that $\eta_\varepsilon(\mathcal{L}) \geq \sqrt{2}\eta_{1/2}$

→ very small ε , larger than necessary BDD radius, too expensive BDD

New idea:

- ▶ find a well-chosen lattice $\mathcal{L} \subset \mathcal{L}' \subset \mathcal{L}/2$ such that $\eta_{\varepsilon'}(\mathcal{L}') \leq \eta_\varepsilon(\mathcal{L})/\sqrt{2}$ for $\varepsilon' \approx \varepsilon$ [ADRS15]
- ▶ run DGS on \mathcal{L}' at $s = \eta_{1/3}(\mathcal{L}) \geq \sqrt{2}\eta_{1/2}(\mathcal{L}')$ [ADRS15]
- ▶ only keep samples in \mathcal{L} (rejection)

DGS sampling: new lemma

- ▶ [ADRS15]: DGS of parameter $s \geq \sqrt{2}\eta_{1/2}(\mathcal{L})$ in time $2^{n/2}$
- ▶ BDD to DGS reduction requires $s = \eta_\varepsilon(\mathcal{L})$ for some $\varepsilon > 0$

Previous work [CCL18]: find ε such that $\eta_\varepsilon(\mathcal{L}) \geq \sqrt{2}\eta_{1/2}$

→ very small ε , larger than necessary BDD radius, too expensive BDD

New idea:

- ▶ find a well-chosen lattice $\mathcal{L} \subset \mathcal{L}' \subset \mathcal{L}/2$ such that $\eta_{\varepsilon'}(\mathcal{L}') \leq \eta_\varepsilon(\mathcal{L})/\sqrt{2}$ for $\varepsilon' \approx \varepsilon$ [ADRS15]
- ▶ run DGS on \mathcal{L}' at $s = \eta_{1/3}(\mathcal{L}) \geq \sqrt{2}\eta_{1/2}(\mathcal{L}')$ [ADRS15]
- ▶ only keep samples in \mathcal{L} (rejection)

Some details:

- ▶ \mathcal{L}' is chosen randomly, works with high probability
- ▶ need that $|\mathcal{L}' / \mathcal{L}| \approx 2^{n/2}$ for $\varepsilon \approx \varepsilon'$
- ▶ rejection: $|\mathcal{L}' / \mathcal{L}| \approx 2^{n/2}$ slowdown, **still better than previous work!**
- ▶ allows to choose $\alpha = 1/3$ for BDD, improved from 0.391 [CCL18]